
COT 6405 Programming Project

Christian Garbin

Contents

1	Introduction	4
2	Longest common subsequence	4
2.1	Definiton	4
2.2	Applications	5
3	Notebook structure	5
4	Algorithm descriptions	6
4.1	Brute-force	6
4.2	Dynamic programming	6
4.3	Hirschberg's linear space	8
4.4	Runtime and space summary	9
5	Planning of experiments	9
5.1	Input size for tests	9
5.2	Runs	10
5.3	Data structures	10
5.4	Input generation	10
5.5	Graphs and tables	10
5.6	Programming language	12
6	Initialization and verification	13
7	Runtime tests and analysis	13
7.1	Constant c calculation	13
7.2	Empirical vs. predicted RT graphs	16
7.3	Graphs for larger input sizes	18
8	Space tests and analysis	20
9	Conclusions	23
9.1	Runtime	23
9.2	Space	24
9.3	Methodology and tools	24
10	Code structure and description	25
10.1	Code structure	25
10.2	How reproducibility is ensured	25
10.3	How time was measured	25
10.4	How memory was measured	25
10.5	Code optimizations	25

11 References	25
12 Appendix	26

Comparison of longest common subsequence (LCS) algorithms

COT 6405 Analysis of Algorithms, Spring 2020

Christian Garbin

1 Introduction

This notebook compares **longest common subsequence** (defined below) algorithms:

- **Brute-force**: generates combinations of subsequences and check if they are common subsequences.
- **Dynamic programming**: takes advantage of common subproblems to not evaluate the same subsequence more than once.
- **Hirschberg's linear space**: a dynamic programming approach, combined with divide-and-conquer, that uses significantly less space than the dynamic programming algorithm.

The comparison measures:

- Runtime efficiency: how long it takes to find a longest common subsequence.
- Space efficiency: how much space is used to find a longest common subsequence.

The code used in the experiments is written in Python 3.x. The code is available in this GitHub repository.

2 Longest common subsequence

2.1 Definition

Given a sequence $X = \langle x_1, x_2, \dots, x_m \rangle$, another sequence Z is a **subsequence** of X if there is a strictly increasing sequence $\langle i_1, i_2, \dots, i_k \rangle$ of indices of X such that for all $j = 1, 2, \dots, k$, we have $x_{i_j} = z_j$ [CLRS01].

For example:

- Given the sequence $X = \langle A, B, C, B, D, A, B \rangle$
- The sequence $Z = \langle B, C, D, B \rangle$ is a subsequence of X , with indices $\langle 2, 3, 5, 7 \rangle$

Given two sequences, a **common subsequence** is a sequence that is common to both sequences. A **longest common subsequence** (LCS) is a maximum-length common subsequence.

For example:

- Given the sequences $X = \langle A, B, C, B, D, A, B \rangle$ and $Y = \langle B, D, C, A, B, A \rangle$

- The sequence $\langle B, C, A \rangle$ is a common subsequence
- The sequence $\langle B, C, B, A \rangle$ is a longest common subsequence
- The sequence $\langle B, D, A, B \rangle$ is another LCS, therefore LCSs do not have to be unique

2.2 Applications

Applications of LCS include determining if two organisms are similar by comparing their DNAs. In this case, "similar" can be determined by the longest common subsequence between the DNAs. The longer the subsequence, the more common the organisms are.

Building on the DNA example, LCS can be used as a form of compression. Using a pre-built suffix tree, an LCS-based solution compressed the human genome from about 3 GB to just over 8 MB [BAF16].

LCS is also used in version control systems to produce the "diff", the minimal amount of additions and deletions that transform the older version of a file into the new version, also known as the "edit distance" [NAV01]. For example, this output of Git's `diff` command between two versions of a file shows deletion and addition of two lines to transform the old version of the file into the new one:

```
1 @@ -33,15 +33,11 @@ class ModifiedFile(BaseCommit):
2     '''Number of lines added to the file in this commit.'''
3     deleted = Column(Integer)
4 -     '''Number of lines deleted from the file in tihs commit.'''
5 +     '''Number of lines deleted from the file in this commit.'''
6 -     commit_id = Column(String)
7 +     commit_id = Column(String, ForeignKey('commit_info.commit_id'))
```

3 Notebook structure

The remainder of this notebook is structured as follows:

- **Algorithms:** describes the algorithms used in the tests.
- **Planning of experiments:** describes the experiments performed, data collected for each one, the analysis performed on them.
- **Initialization and verification:** initializes the notebook (import Python modules, set important environment values) and verifies that the algorithms are working before using them.
- **Measurement and analysis:** run the tests to collect the metrics and analyzes those metrics.
- **Code:** documents relevant pieces of the code and other technical aspects found during the development and execution of the tests.

4 Algorithm descriptions

Three algorithms will be analyzed:

- Brute-force
- Dynamic programming
- Hirschberg's linear space

In all sections below, m is the length of the longest sequence and n is the length of the shortest sequence.

4.1 Brute-force

Brute-force is the simplest LCS algorithm: generate subsequences of the smaller sequence and check if they are also a subsequence of the larger sequence.

The pseudocode for the algorithm is:

```

1  LCS_BRUTE_FORCE(X, Y)
2      // Pick the shortest sequence to generate subsequences
3      short_seq = shortest_of(X, Y)
4      long_seq = longest_of(X, Y)
5
6      // Try all subsequences of the shortest sequence
7      for i = length(short_seq) to 1
8          // Try all subsequences of length i, one at a time
9          while "there are subsequences of length i to try":
10             subseq = next_subsequence(short_seq, i)
11             if is_subsequence(subseq, long_seq)
12                 return subseq
13
14     // Could not find a subsequence
15     return [] // empty sequence

```

Runtime analysis: the algorithm selects the smaller sequence to generate combinations to test. Since there are 2^n combinations of subsequences that can be generated from the subsequence, and each one has to be tested against the larger sequence, the runtime is $m \times 2^n$. In most cases, the second term is much larger than m , making it a $O(2^n)$ algorithm.

Space analysis: A naive implementation would generate all combinations of the smaller sequence ahead time, using $O(2^n)$ space. An optimized implementation, as shown above, generates one combination of the smaller sequence at a time, using $O(n)$ space.

4.2 Dynamic programming

Dynamic programming makes use of the optimal substructure of the LCS, solving smaller subproblems only once, combining the solutions.

The pseudocode is shown below. It has two parts: first two matrices are constructed to determine the LCS length and how to construct it (a series of "moves"), then the LCS is extracted by going through the moves matrix.

```

1  LCS_LENGTH(X, Y)
2      m = length(X)
3      n = length(Y)
4
5      // c is an m x n matrix with the top row and
6      // left column initialized to zero
7      c = matrix(m, n)
8      for i = 1 to m
9          c[i, 0] = 0
10     for j = 1 to n
11         c[0, j] = 0
12
13     // b is an m x n empty matrix that will hold the
14     // movements to build the LCS
15     b = matrix(m, n)
16
17     for i = 1 to m
18         for j = 1 to n
19             if X[i] == Y[j]
20                 c[i,j] = c[i-1,j-1] + 1
21                 b[i,j] = "diagonal"
22             else if c[i-1,j] >= c[i, j-1]
23                 c[i,j] = c[i-1,j]
24                 b[i,j] = "up"
25             else
26                 c[i,j] = c[i,j-1]
27                 b[i,j] = "down"
28
29     // c[m,n] has the LCS length and b has the
30     // sequences of moves to extract the LCS
31     return c, b
32
33 EXTRACT_LCS(b, X, i, j)
34     lcs = empty_list()
35
36     while i > 0 and j > 0
37         move = b[i,j]
38         if move == "diagonal"
39             lcs = lcs + X[i]
40             i = i - 1
41             j = j - 1
42         else if move == "up"
43             i = i - 1
44         else // "down"
45             j = j - 1
46
47     // The LCS was built from the bottom up,
48     // need revert it before returning
49     return reverse(lcs)
50
51 LCS(X, Y)
52     b, c = LCS_LENGTH(X, Y)
53     lcs = EXTRACT_LCS(b, X, length(X), length(Y))
54     return lcs

```

Runtime analysis: the `LCS_LENGTH` part of the algorithm is $O(m \times n)$, from its two nested loops. The `EXTRACT_LCS` is $O(m + n)$. For large values of m and n , `LCS_LENGTH` dominates the runtime, making the algorithm overall $O(m \times n)$.

Space analysis: the $m \times n$ matrix in `LCS_LENGTH` is responsible for the space the algorithm needs, thus the space is $O(m \times n)$.

4.3 Hirschberg's linear space

Hirschberg's linear space algorithm [HIR75] is a dynamic programming approach that uses divide-and-conquer. As the name indicates, it makes efficient use of space.

The pseudocode is shown below. It has two parts: a *scoring* (also called *cost*) function to help decide where to divide the current subsequence being analyzed, and the function that divides-and-conquers, based on that score.

```

1  SCORE(X, Y)
2    m = length(X)
3    n = length(Y)
4
5    // A list of of scores, initialized with n zeros
6    scores = list(0 * n)
7
8    for i = 1 to m
9      prev_score = scores
10     for j = 1 to n
11       if X[i] == Y[j]
12         scores[j+1] = prev_scores[j] + 1
13       else
14         scores[j+1] = max(scores[j], prev_scores[j + 1])
15
16    return scores
17
18  LCS(X, Y)
19    m = length(X)
20    n = length(Y)
21    if m == 0
22      // Got to the end of the sequence
23      return []
24    else if m == 1
25      // Last character, check if it is in subsequence
26      if X[1] is in Y
27        return X[1]
28      else
29        return []
30    else
31      // Find where to split the current sequences
32
33      // X is split in the middle
34      i = m / 2
35      XB = X[1:i]
36      XE = X[i+1:m]
37
38      // Y is split based on the scores
39      cost_top_left = SCORE(XB, Y)
40      cost_bottom_right = SCORE(reverse(XE), reverse(Y))
41      cost = cost_top_left + reverse(cost_bottom_right)
42
43      k = index of max(cost) // argmax(cost)
44      YB = Y[1:k]
45      YE = Y[k+1:n]
46
47      // Solve for each part of the split sequences
48      return LCS(XB, YB) + LCS(XE, YE)

```

Runtime analysis: in each step, the sequences under examination are split into two subsequences at $m/2$ and a q based on a cost factor (this is the secret cause of the algorithm). It can be shown that the recurrence is $O(mn)$ [KT05] [FAG16].

Space analysis: a naive implementation of the algorithm creates copies of the sequences as it splits them during the recursive calls, using space $O(m+n)$. An implmentation that passes the

original sequences around and uses indices to logically split them (without creating copies), uses space $O(\min(m, n))$ (in the `SCORE` fuction).

4.4 Runtime and space summary

The following table summarizes the runtime and space characteristics of the algorithms.

Algorithm	Runtime	Space
Brute-force	$O(2^n)$	$O(n)$
Dynamic programming	$O(m \times n)$	$O(m \times n)$
Hirschberg	$O(m \times n)$	$O(\min(m, n))$

5 Planning of experiments

The experiments compare the runtime and space of the brute-force, dynamic programming recursive, and Hirschberg's linear space algorithms.

To illustrate the algorithms in a typical application, the tests will use two strings that resemble DNA sequences (a combination of the letters `A`, `C`, `G`, and `T`) and will find an LCS for them. To emulate the computational biology case of searching for a common substring between two DNA strands, in each case we will search for a string that is one-tenth of the larger strings, illustrated in the table below.

5.1 Input size for tests

Strings of three sizes will be used, small, medium, large. For each size, the same strings will be used with all algorithms, to keep the comparison consistent.

Test size	The DNA strain (X) This is the m in RT and space analysis	The possible common sequence (Y) This is the n in RT and memory analysis
Small	1,000	100
Medium	10,000	1,000
Large	100,000	10,000

Table 1 - Size of strings to test and how they map to the m and n of the RT and memory analysis.

5.2 Runs

Each algorithm will be executed ten times ($k=10$) for each string size to remove variations in the environment. The average of these runs will be used as the final number for the algorithm.

Two values will be measured in each run, running time (RT) and memory (space) usage:

- Time: measured with Python's `time` package.
- Memory: measured with Python's `memory_profiler` package.

Details of how measurements were conducted are documented in the code section, later in this document.

5.3 Data structures

Experimental data, the strings, will be stored in the standard data structures for string representation, usually mapped to a constant-time access continuous array in programming languages. Auxiliary data structures to keep track of intermediate results will be kept either in Python arrays or NumPy arrays, whichever is more performant for a specific piece of code.

5.4 Input generation

Strings for the tests will be generated using a pseudo-random number generation initialized with a seed, to ensure the repeatability of the experiments (the same sequence is generated every time). The strings will be generated only once, before each algorithm is executed, to ensure that the results can be compared with each other.

5.5 Graphs and tables

For each algorithm, two tables will be filled in:

- RT analysis: theoretical vs. empirical RT.
- Memory usage: theoretical vs. empirical memory usage.

The following table illustrates the RT analysis for the brute-force algorithm.

Test size	Theoretical complexity	Empirical RT (ms)	Ratio (empirical RT / theoretical complexity)		
			Predicted RT	% error	
Small $m=1,000$, $n=100$	$2^n = 2^{100}$				

Test size	Theoretical complexity	Empirical RT (ms)	Ratio (empirical RT / theoretical complexity)	Predicted RT	% error
Medium m=10,000, n=1,000	$2^n = 2^{1000}$				
Large m=100,000, n=10,000	$2^n = 2^{10000}$				

Table 2 - RT analysis table example, using the brute-force algorithm as illustration

The table columns are computed as follows:

- *Ratio* measures the ratio between the empirical and the theoretical complexity. Its value is always > 0 . It is used to calculate the constant c .
 - The constant c measures the overhead of the steps (computer instructions) that are outside the main loops (or recursive calls) of the algorithms. It is determined as the maximum of the ratio values (with outliers discarded), i.e. the maximum value of the *Ratio* column.
- *Predicted RT* is computed as the constant c times the *Theoretical complexity*.
- *% error* measures the discrepancy between the predicted and the empirical time: $(Empirical\ RT - Predicted\ RT) / Empirical\ RT * 100$.

The following table illustrates memory usage for the brute-force algorithm.

Test size	Theoretical memory usage (KiB)	Empirical memory usage (KiB)	% error
Smallm=1,000, n=100	$n / 1024 = 0.98$		
Mediumm=10,000, n=1,000	9.8		
Large m=100,000, n=10,000	97.7		

Table 3 - Memory usage analysis table example, using the brute-force algorithm as illustration

The table columns are computed as follows:

- *Theoretical memory usage* is the number of characters needed, times 1 byte per character,

divided by 1,024 to transform to KiB.

- "1 byte per character" comes from the variable-length encoding of strings in Python using the CPython environment. Because we are representing DNA strands, we are using only ASCII characters, which are represented as 1 byte [CPY20] [GOL20].
- % error measures the discrepancy between the theoretical and empirical memory usage: $\text{Empirical memory usage} / \text{Predicted memory usage} / \text{Empirical memory usage} * 100$.

For the other algorithms, the theoretical values will be adjusted as follows:

- Dynamic programming:
 - Theoretical RT = $m \times n$
 - Theoretical memory usage = $m \times n$
- Hirschberg's linear space algorithm
 - Theoretical RT = $m \times n$
 - Theoretical memory usage = $m + n$

Once the tables are filled in, two sets of graphs will be created:

- **Algorithm comparison:** this set of graphs compares the empirical runtime and memory usage of the algorithms. There will be one runtime and one memory usage plot for all algorithms (two graphs).
- **Theoretical vs. empirical results:** this set of graphs compares the empirical runtime and memory usage of each algorithm. There will be one set of plots for each algorithm and each size, for a total of 9 graphs (3 algorithms, 3 sizes).

The following graphs will be generated for algorithm comparison:

- RT comparison: a horizontal bar graph with the runtime in ms (horizontal axis) for each algorithm, grouped by the input size. A horizontal bar graph will be used because of the expected large values for large input sizes. This representation makes better use of space.
- Memory comparison: similar to the graph above, using memory usage as the horizontal axis.

The following graphs will be generated for each algorithm for the theoretical vs. empirical results:

- Theoretical vs. empirical runtime in ms for each input size
- Theoretical vs. empirical memory usage in KiB for each input size

5.6 Programming language

The experiments use Python 3.x in a Jupyter Notebook environment.

6 Initialization and verification

Load commonly-used modules.

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 pd.set_option('precision', 3)
```

Check that the algorithms work by testing them against controlled input.

There are three part to the tests:

1. Automated tests that check against well-defined inputs. They are meant to be easy to debug, in case an algorithm fails.
2. Tests with longer inputs that simular DNA strands. They test more realistic scenarios, but still short enough to run fast.
3. A visual check, by printing the aligned subsequence. They guard against the test code itself having a failure that generates false positives.

```
1 import utils.lcs_test
2
3 lcs_test.test(visualize=True)
```

```
1 All basic tests passed
2 All DNA tests passed
3 Visual inspection:
4               lcs_brute_force:  CACATTGCCTGGATAGGGGCTAGGATCGAG
5               lcs_brute_force:  ....TT.....AT.GG....AGGA.CGA. size=13
6       lcs_dynamic_programming_matrix_numpy:  ....TT.....AT.GG....AGGA.CGA. size=13
7               lcs_hirschberg_numpy:  ....TT.....AT.GG....AGGA.CGA. size=13
```

7 Runtime tests and analysis

To illustrate a real-life scenario, the code checks if a DNA strand is part of a larger DNA sequence [WIK20a].

```
1 # Force reload because this piece of code frequently changes
2 import importlib
3 import metrics as m
4 importlib.reload(m);
```

7.1 Constant c calculation

This section calculates the constant c for each algorithm. This is the constant that accounts for instructions that are not in the loops. For example, given a runtime $O(n^2)$, the constant c allows us to write the runtime more precisely as $c \times n^2$.

The constant is also affected by the language and compiler or interpreter used. In this notebook it is calculated with dynamic analysis: run the code and calculate the overhead.

Each algorithm was executed once with different input sizes, with the pairs representing the length of X and Y, respectively: (1,000, 100), (2,000, 200), (3,000, 300), (4,000, 300), (4,000, 500), (4,000, 1,000).

Once the algorithms are run, we calculate c as $\max(rt_1, rt_2, \dots, rt_n)$, excluding outliers where applicable.

```
1 rt_results_raw, rt_results_summary = m.runtime(m.seq_phase1, verbose=1,
2       file='runtime-phase1')
```

```
1 Loading from file
```

We now have two Pandas DataFrames:

- `rt_results_raw`: results from all 10 executions of each algorithms and each input size.
- `rt_results_summary`: average of the executions for each each algorithm and input size.

Using these DataFrames, the following sections calculate three constants, one for each algorithm. They are stored in the following variables:

- `c_bf`: the constant for the brute-force algorithm.
- `c_dp`: the constant for the dynamic programming algorithm.
- `c_h`: the constant for the Hirschberg linear space algorithm.

The dataframes contain metrics for all algorithms. The sections below filter the dataframe for the algorithm analyzed as needed.

7.1.1 Constant c for brute-force

```
1 rt_bf, c_bf = m.add_runtime_analysis(rt_results_summary, m.ALG_BRUTE_FORCE)
2 print('c_bf={}'.format(c_bf))
3 display(rt_bf)
```

```
1 c_bf=5.206481974415714e-32
```

	Algorithm	Sequence size	Subsequence size	Empirical RT (ms)	\
0	Brute-force	1000	100	0.066	
1	Brute-force	2000	200	0.173	
2	Brute-force	3000	300	0.206	
3	Brute-force	4000	300	0.256	
4	Brute-force	4000	500	0.495	
5	Brute-force	4000	1000	156.909	
6	Brute-force	5000	900	0.601	
7	Brute-force	5000	1000	0.658	
8	Brute-force	5000	1200	0.760	
	Theoretical complexity	Ratio	Predicted RT	% error	
0	1.268e+30	5.206e-32	6.600e-02	0.000e+00	
1	1.607e+60	1.077e-61	8.366e+28	-4.836e+31	
2	2.037e+90	1.011e-91	1.061e+59	-5.148e+61	
3	2.037e+90	1.257e-91	1.061e+59	-4.143e+61	
4	3.273e+150	1.512e-151	1.704e+119	-3.443e+121	
5	1.072e+301	1.464e-299	5.579e+269	-3.555e+269	
6	8.453e+270	7.110e-272	4.401e+239	-7.323e+241	
7	1.072e+301	6.141e-302	5.579e+269	-8.478e+271	
8	inf	0.000e+00	inf	-inf	

We can see in the table a very large error for most tests of the brute-force algorithm. This is caused by the "luck factor" of this algorithm: if we are lucky and generate a combination early on that happens to be a common subsequence, the algorithm terminates quickly. The probability of generating a common subsequence is high in this case because of the reduced amount of possible combinations we have when using only the four letter of a DNA sequence.

Given the "luck factor" of this case, analyzing the predicted vs. empirical runtime will not be insightful.

7.1.2 Constant c for dynamic programming

```
1 rt_dp, c_dp = m.add_runtime_analysis(rt_results_summary, m.ALG_DYNAMIC_PROGRAMMING)
2 print('c_dp={}'.format(c_dp))
3 display(rt_dp)
```

```
1 c_dp=0.0001627695555555558
```

	Algorithm	Sequence size	Subsequence size	Empirical RT (ms)	\
2	9 Dynamic programming	1000	100	15.257	
3	10 Dynamic programming	2000	200	62.163	
4	11 Dynamic programming	3000	300	142.757	
5	12 Dynamic programming	4000	300	182.718	
6	13 Dynamic programming	4000	500	308.720	
7	14 Dynamic programming	4000	1000	612.753	
8	15 Dynamic programming	5000	900	732.463	
9	16 Dynamic programming	5000	1000	777.084	
10	17 Dynamic programming	5000	1200	938.708	
11					
	Theoretical complexity	Ratio	Predicted RT	% error	
13	9 100000	1.526e-04	16.277	-6.685	
14	10 400000	1.554e-04	65.108	-4.737	
15	11 900000	1.586e-04	146.493	-2.617	
16	12 1200000	1.523e-04	195.323	-6.899	
17	13 2000000	1.544e-04	325.539	-5.448	
18	14 4000000	1.532e-04	651.078	-6.255	
19	15 4500000	1.628e-04	732.463	0.000	
20	16 5000000	1.554e-04	813.848	-4.731	
21	17 6000000	1.565e-04	976.617	-4.038	

Since all results are within a small margin of error, none of them will be discarded for the calculation.

7.1.3 Constant c for Hirschberg's linear space

```
1 rt_h, c_h = m.add_runtime_analysis(rt_results_summary, m.ALG_HIRSCHBERG)
2 print('c_h={}'.format(c_h))
3 display(rt_h)
```

```
1 c_h=0.00034574999999999363
```

	Algorithm	Sequence size	Subsequence size	Empirical RT (ms)	\
2	18 Hirschberg	1000	100	34.575	
3	19 Hirschberg	2000	200	135.058	
4	20 Hirschberg	3000	300	290.896	
5	21 Hirschberg	4000	300	387.301	
6	22 Hirschberg	4000	500	621.257	
7	23 Hirschberg	4000	1000	1226.753	

8	24	Hirschberg	5000	900	1399.517
9	25	Hirschberg	5000	1000	1542.752
10	26	Hirschberg	5000	1200	1846.447
11					
12		Theoretical complexity	Ratio	Predicted RT	% error
13	18	100000	3.457e-04	34.575	0.000
14	19	400000	3.376e-04	138.300	-2.400
15	20	900000	3.232e-04	311.175	-6.971
16	21	1200000	3.228e-04	414.900	-7.126
17	22	2000000	3.106e-04	691.500	-11.307
18	23	4000000	3.067e-04	1383.000	-12.737
19	24	4500000	3.110e-04	1555.875	-11.172
20	25	5000000	3.086e-04	1728.750	-12.056
21	26	6000000	3.077e-04	2074.500	-12.351

Some of the results are off by more than 10%, but not by much, so we will keep all of them for the calculation.

7.2 Empirical vs. predicted RT graphs

The graphs below compare the empirical runtime with the predicted runtime using the smaller input sizes. The predicted runtime is calculated as $c \times \text{theoretical complexity}$.

Each of the graphs below shows the empirical vs. predicted runtime for an algorithm. The empirical runtime is the actual running time of the algorithms, while the predicted runtime is calculated as $c \times$ the theoretical complexity.

```

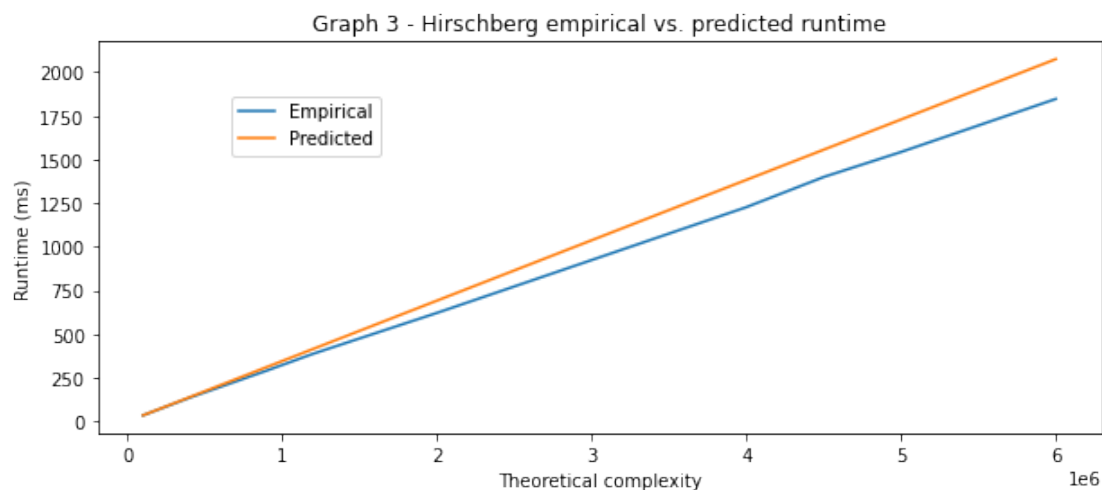


1 graph_number = 0
2 def plot_empirical_predicted_rt(data, title):
3     '''Auxiliary function to plot the graphs consistently for all algorithms.'''
4     global graph_number
5     graph_number += 1
6     fig = plt.figure(figsize=(10,4))
7     sns.lineplot(x=m.DF_THEORETICAL_COMPLEXITY, y=m.DF_EMPIRICAL_RT, data=data)
8     sns.lineplot(x=m.DF_THEORETICAL_COMPLEXITY, y=m.DF_PREDICTED_RT, data=data)
9     plt.title('Graph {} - {} empirical vs. predicted runtime'.format(
10         graph_number, title))
11     plt.ylabel("Runtime (ms)")
12     fig.legend(['Empirical', 'Predicted'], bbox_to_anchor=(0.3, 0.8))
13     plt.show()

```

```

1 for alg in (rt_bf, rt_dp, rt_h):
2     plot_empirical_predicted_rt(alg, alg.iloc[0][m.DF_ALGORITHM])

```

From these graphs we can see that:

1. The brute-force graph shows again the "luck" effect, where the runtime is low because a subsequence just happens to be found early on. It could have gone the other way, as it in fact did in one of the experiments, where after running for eight hours, it still did not find a common subsequence for $m = 5,000$ and $n = 2,000$.

- Both the dynamic programming and the Hirschberg algorithms track closely to their predicted runtime.
- The Hirschberg algorithm's error is larger for larger input sizes, likely because the nature of the implementation and the programming language. It is a recursive algorithm and Python does not support tail recursion optimization [WIK20b] [ROS09]. Thus the loops have the extra cost of function calls that the dynamic programming algorithm does not have. The cost of the function calls is relatively more expensive for small input size, resulting in overestimating c . An improvement for this case could be to do a more rigorous outlier elimination when calculating c .

7.3 Graphs for larger input sizes


The graphs below compare the empirical runtime with the predicted runtime using the larger input sizes. As before, the predicted runtime is calculated as $c \times \text{theoretical complexity}$.


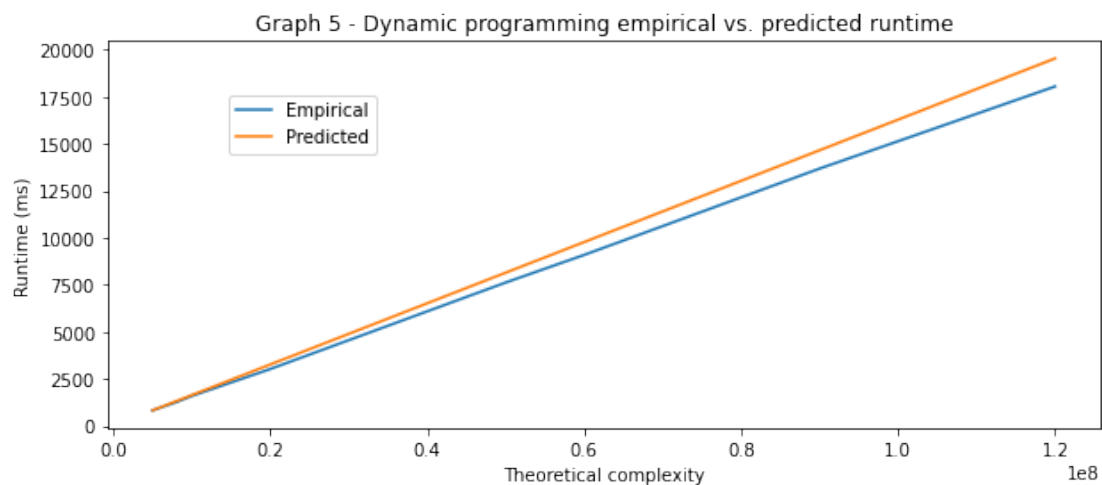
The graphs show the same trend lines as the graphs for the smaller input sizes.

```
1 rt_results_raw, rt_results_summary = m.runtime(m.seq_phase2, verbose=1,
2       file='runtime-phase2')
```

```
1 Loading from file
```

```
1 for alg in (m.ALG_BRUTE_FORCE, m.ALG_DYNAMIC_PROGRAMMING, m.ALG_HIRSCHBERG):
2     rt, _ = m.add_runtime_analysis(rt_results_summary, alg)
3     plot_empirical_predicted_rt(rt, rt.iloc[0][m.DF_ALGORITHM])
```






Graphing the algorithms separately does not clearly show how their runtime compare with each other. The graph below shows the dynamic programming and the Hirschberg algorithm with the same scale for the vertical axis (the brute-force algorithm is not shown because its complexity is in another scale entirely, therefore not comparable with the other two algorithms).

With this graph we can see that although they have the same growth rate (both have $O(mn)$ complexity), the Hirschberg algorithm has a higher slope, resulting from its higher c constant.

```
1 dp_summary, _ = m.add_runtime_analysis(rt_results_summary, m.
    ALG_DYNAMIC_PROGRAMMING)
2 h_summary, _ = m.add_runtime_analysis(rt_results_summary, m.ALG_HIRSCHBERG)
3 summary = pd.concat([dp_summary, h_summary])
```

```
1 graph_number += 1
2 g = sns.FacetGrid(summary, col=m.DF_ALGORITHM, height=4, aspect=1.5)
3 g.map(plt.plot, m.DF_THEORETICAL_COMPLEXITY, m.DF_PREDICTED_RT,
4     color='darkorange')
5 g.map(plt.plot, m.DF_THEORETICAL_COMPLEXITY, m.DF_EMPIRICAL_RT,
6     color='royalblue')
7 g.set_ylabels('Runtime (ms)').fig.subplots_adjust(top=0.8)
8 g.fig.suptitle('Graph {} - Dynamic programming compared to Hirschberg'.format(
9     graph_number))
10 plt.show()
```



8 Space tests and analysis

```
1 mem_results_raw, mem_results_summary = m.memory(m.seq_phase2, verbose=1,
2         file='memory-phase2')
3 mem_results_summary = m.add_memory_analysis(mem_results_summary)
```

```
1 Loading from file
```

We now have two Pandas DataFrames:

- `mem_results_raw`: results from all 10 executions of each algorithms and each input size.
- `mem_results_summary`: average of the executions for each each algorithm and input size.

The DataFrames contain metrics for all algorithms. Sections below filters the row for the algorithm analyzed in that section.

The tables below show the empirical and predicted space in MiB. Using these values, an error column is also shown.

Observations from these tables:

1. Measuring small amounts of memory used in short period of times, as used by the brute-force algorithm in all cases and by the other algorithms with smaller input sizes, is unreliable. There is a discussion about the method used later in the notebook, in the code section. Based on these observations, the results from the brute-force algorithm will not be used in further analysis.
2. As the runtime increases and the amount of memory used grows larger, the measurements become more reliable.

```
1 for alg in [m.ALG_BRUTE_FORCE, m.ALG_DYNAMIC_PROGRAMMING, m.ALG_HIRSCHBERG]:
2     df = mem_results_summary[mem_results_summary[m.DF_ALGORITHM] == alg]
3     print('\n{} memory analysis'.format(alg))
4     display(pd.pivot_table(df, values=[m.DF_EMPIRICAL_SPACE, m.DF_PREDICTED_SPACE,
5         m.DF_ERROR],
6         index=[m.DF_SEQ_SIZE, m.DF_SUBSEQ_SIZE]))
```

```
1 Brute-force memory analysis
```

```
1                                     % error  Empirical space (MiB)  \
2 Sequence size Subsequence size
```

3	10000	500	-inf	0.000
4		800	-inf	0.000
5		1000	-inf	0.000
6	20000	1000	-inf	0.000
7		2000	-inf	0.000
8		2500	38.965	0.004
9	30000	2000	-inf	0.000
10		3000	75.586	0.012
11		4000	87.793	0.031
12				
13			Predicted space (MiB)	
14	Sequence size	Subsequence size		
15	10000	500	4.768e-04	
16		800	7.629e-04	
17		1000	9.537e-04	
18	20000	1000	9.537e-04	
19		2000	1.907e-03	
20		2500	2.384e-03	
21	30000	2000	1.907e-03	
22		3000	2.861e-03	
23		4000	3.815e-03	

1 Dynamic programming memory analysis

1			% error	Empirical space (MiB)	\
2	Sequence size	Subsequence size			
3	10000	500	-162660.417	0.012	
4		800	0.134	30.559	
5		1000	0.167	38.211	
6	20000	1000	0.101	76.371	
7		2000	0.058	152.676	
8		2500	0.047	190.824	
9	30000	2000	0.057	229.012	
10		3000	0.037	343.449	
11		4000	0.029	457.895	
12					
13			Predicted space (MiB)		
14	Sequence size	Subsequence size			
15	10000	500	19.073		
16		800	30.518		
17		1000	38.147		
18	20000	1000	76.294		
19		2000	152.588		
20		2500	190.735		
21	30000	2000	228.882		
22		3000	343.323		
23		4000	457.764		

1 Hirschberg memory analysis

1			% error	Empirical space (MiB)	\
2	Sequence size	Subsequence size			
3	10000	500	-inf	0.000	
4		800	92.188	0.039	
5		1000	91.862	0.047	
6	20000	1000	91.862	0.047	
7		2000	91.862	0.094	
8		2500	91.862	0.117	
9	30000	2000	86.979	0.059	
10		3000	97.287	0.422	
11		4000	97.543	0.621	
12					
13			Predicted space (MiB)		
14	Sequence size	Subsequence size			
15	10000	500	0.002		
16		800	0.003		
17		1000	0.004		
18	20000	1000	0.004		
19		2000	0.008		

20		2500	0.010
21	30000	2000	0.008
22		3000	0.011
23		4000	0.015

Memory analysis for the dynamic programming and Hirschberg's linear space is more interesting.

Before going into the analysis, a review of some implementation details that affect the space characteristics of the algorithms:

1. The classic dynamic programming algorithm uses two $m \times n$ matrices, one for the length and another for the direction of the moves. Space-optimized implementations combine these matrices into one, reserving bits in each cell for the length and direction. Such an implementation was used here. It is discussed in the code section of the notebook.
2. Hirschberg's traditional implementation is recursive. In languages that do not support tail recursion, such as Python, each recursion creates a new stack frame. Thus, some of the memory used by the algorithm is in the form of the stack frames, in addition to the arrays it needs for the algorithm itself.
3. Also for Hirschberg's, a simplistic implementation of the algorithm creates copies of the sequences as it finds where to split them. A space-optimized algorithm uses indices into the original sequence to avoid creating copies. Such an implementation was used here. It is discussed in the code section of the notebook.

With that in mind, from the tables we observe that:

1. The dynamic programming algorithm tracks closely to the predicted space. This is due to two factors. First, it uses a large amount of memory, which seems to favor this particular method of measuring memory utilization (discussion in the code section). And second, it uses loops, as opposed to recursion, which does not create the overhead of stack frames (in languages without tail recursion).
2. The Hirschberg's linear space algorithm has a large error. This error is likely caused by it being a recursive algorithm. Some of the memory measured during the execution comes from the stack frames created in each recursion.

Finally, the most important observation is the difference between the dynamic programming and Hirschberg's linear space algorithm. As expected, dynamic programming uses significantly more memory.

On the other hand, in this implementation, dynamic programming is twice as fast as Hirschberg's. In another programming language, with better support for tail recursion, their performance may be comparable.

In most real-life applications, with large input sizes, a fine-tuned implementation of Hirschberg's linear space is the preferable option. In some cases, for very large input sizes, it may be the only feasible option.

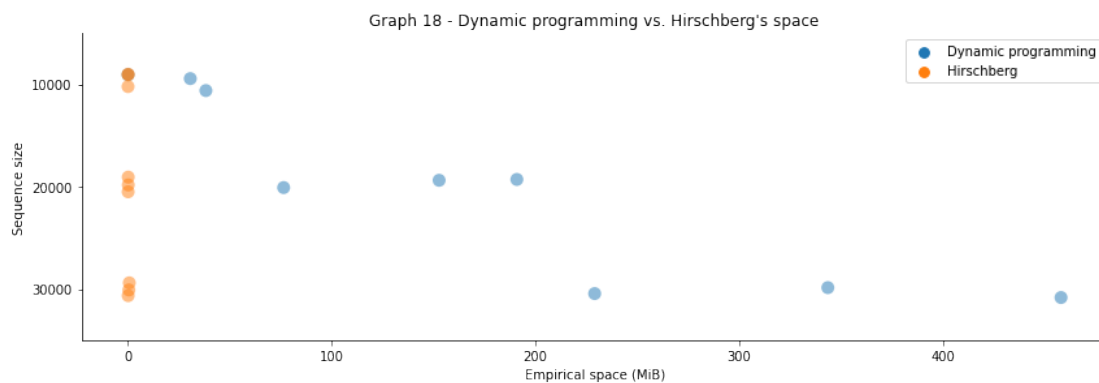
To illustrate the dramatic difference in space, the graph below shows the empirical space used by the two algorithms side-by-side, for each input size. The horizontal axis is the empirical

space. Each dot represent a subsequence for a sequence size (vertical axis). We can see that Hirschberg's linear space algorithm (orange) never goes past one MiB, while the dynamic programming algorithm (blue) escalates quickly, going past hundreds of MiBs, reaching almost half a gigabyte for the largest combination of sequences.

```

1 graph_number += 1
2 df = mem_results_summary[(mem_results_summary[m.DF_ALGORITHM] == m.
    ALG_DYNAMIC_PROGRAMMING) | (mem_results_summary[m.DF_ALGORITHM] == m.
    ALG_HIRSCHBERG)]
3 sns.catplot(y=m.DF_SEQ_SIZE, x=m.DF_EMPIRICAL_SPACE, hue=m.DF_ALGORITHM, data=df,
    orient='h', height=4, aspect=3, legend=False, s=10, alpha=0.5)
4 plt.title('Graph {} - Dynamic programming vs. Hirschberg's space'.format(
5     graph_number))
6 plt.legend(loc='upper right')
7 plt.show()

```



9 Conclusions

9.1 Runtime

From the runtime experiments we can conclude that:

1. The brute-force algorithm is basically a coin flip, where the odds for "it will run quickly" are much slower than the odds for "it will run very, very slowly". It may find an LCS in extremely fast times (under 100 ms), or it may take hours and still not find an LCS (as it happened once during the research for this project). However, in some applications, where the probability of finding a common subsequence is high (a combination of small input sizes and small set of characters to pick from), running the brute-force algorithm in parallel with a dynamic programming algorithm may be worthwhile. When it pays off, it pays off big.
2. The dynamic programming and Hirschberg's algorithm have stable runtime characteristics, that is given an input size, it is easy to calculate how long it will take to find an LCS. Predictability is a desirable characteristic in real-life applications.

9.2 Space

1. The pseudocode for the algorithms are, as a general rule, not a good example of memory efficiency. For example, the pseudocode for the dynamic programming algorithm usually shows two tables, one for the length and another for the moves. While this approach makes the pseudocode easy to understand, it is not an efficient utilization of space. To make the most space-intensive algorithms work, careful use of memory management techniques is needed.
2. The difference in space utilization between the dynamic programming and the Hirschberg's linear space algorithms is astounding when they are put side by side. It shows one of them barely needing one MiB, while the other reaching half of a gigabyte to perform the same work. Hats off to Dr. Hirschberg.

9.3 Methodology and tools

1. Getting the algorithms to run fast and use a reasonable amount of memory requires knowledge of the particular environment (e.g a Python environment, compared to a C++ environment). The naive implementation, one that follows the pseudocode from textbooks closely, is usually slow, uses too much memory, or both.
2. Python is a good choice for experimentation, but not for performance. For example, for the brute-force algorithm Python's `itertools` is a fast and efficient way to generate combinations. On the other hand, the lack of tail recursion handicaps algorithms that make use of recursion, such as Hirschberg's. This introduces potentially artificial differences between the algorithms, i.e. differences that are caused by the environment, not necessarily by fundamental differences in the algorithms.
3. Jupyter speeds up the "experiment, evaluate" cycle greatly. It also makes the process repeatable and transparent by exposing the code used for analysis, facilitating peer review of the methods and assumptions used in the work.
4. Measuring memory usage was surprisingly hard, especially when the usage was on the low side (less than one MiB). Part of the reason is the environment (Python, which is a mixture of reference counting and garbage collection) and tools (`memory_profile`). If I had to do this again, I would write these pieces of the code in C or C++, as written in another point, and spawn individual processes for each test case, measuring memory usage with operating system tools, instead of language modules/libraries.

To summarize, if I had to do it again, I would have written the algorithms in C or C++ to take advantage of their performance and write the results into a file, then use Python to read the file and perform the analysis in a Jupyter notebook. This approach would combine the best of both worlds.

10 Code structure and description

This section highlights pieces of the code that are significant for the experiments.

10.1 Code structure

10.2 How reproducibility is ensured

10.3 How time was measured

10.4 How memory was measured

10.5 Code optimizations

11 References

[BAF16] Beal, R., Afrin, T., Farheen, A. et al. 2016. *A new algorithm for “the LCS problem” with application in compressing genome resequencing data*. BMC Genomics 17, 544 (2016). <https://doi.org/10.1186/s12864-016-2793-0>, accessed 2020-03-16.

[CLRS01] Cormen, T., Leiserson, C., Rivest, R., and Stein, C. 2001. *Introduction to Algorithms*, 2nd edition. MIT Press.

[CPY20] *The CPython source code*. <https://github.com/python/cpython/blob/v3.6.3/Include/unicodeobject.h#L202>, accessed 2020-03-18.

[FAG16] Fagerberg, R. 2016. *Dynamic Programming: Hirschberg’s Trick*. <https://imada.sdu.dk/~rolf/Edu/DM823/E16/Hirschberg.pdf>, accessed 2020-04-25.

[GOL20] Golubin, A. *How Python saves memory when storing strings*. <https://rushter.com/blog/python-strings-and-memory/>, accessed 2020-03-18.

[HIR75] Hirschberg, D. *A linear space algorithm for computing maximal common subsequences*. Commun. ACM 18 (1975): 341-343.

[KT05] Kleinberg, J. and Tardos, E. 2005. *Algorithm Design*. Addison-Wesley Longman Publishing Co., Inc., USA.

[NAV01] Navarro, G. 2001. *A guided tour to approximate string matching*. ACM Computing Surveys. <http://users.csc.calpoly.edu/~dekhtyar/570-Fall2011/papers/navarro-approximate.pdf>, accessed 2020-03-27.

[ROS09] van Rossum, G. 2009. *Tail recursion elimination*. <http://neopythonic.blogspot.com/2009/04/tail-recursion-elimination.html>, accessed 2020-04-28.

[WIK20a] Wikipedia. *Subsequence* entry, *Applications* section. <https://en.wikipedia.org/wiki/Subsequence#Applications>, accessed 2020-03-18.

[WIK20b] Wikipedia. *Tail call*. https://en.wikipedia.org/wiki/Tail_call, accessed 2020-04-28.

References used in the code are annotated directly in the code.

12 Appendix

Show full data from the experiments here
